

[◀ Return to "Deep Learning" in the classroom](#)[DISCUSS ON STUDENT HUB](#)

Generate TV Scripts

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

Good job overall generating your own Simpsons TV script using Recurrent Neural Networks!

You put a lot of effort into this project!

Congratulations and good luck with your Nanodegree!

Required Files and Tests

The project submission contains the project notebook, called "dlnd_tv_script_generation.ipynb".

All the correct files are included!

All the unit tests in project have passed.

Impressive work getting all the unit tests to pass!

Preprocessing

The function `create_lookup_tables` create two dictionaries:

- Dictionary to go from the words to an id, we'll call `vocab_to_int`
- Dictionary to go from the id to word, we'll call `int_to_vocab`

The function `create_lookup_tables` return these dictionaries in the a tuple (`vocab_to_int`, `int_to_vocab`)

The function `token_lookup` returns a dict that can correctly tokenizes the provided symbols.

Build the Neural Network

Implemented the `get_inputs` function to create TF Placeholders for the Neural Network with the following placeholders:

- Input text placeholder named "input" using the TF Placeholder name parameter.
- Targets placeholder
- Learning Rate placeholder

The `get_inputs` function return the placeholders in the following the tuple (Input, Targets, LearningRate)

The `get_init_cell` function does the following:

- Stacks one or more BasicLSTMCells in a MultiRNNCell using the RNN size `rnn_size`.
- Initializes Cell State using the MultiRNNCell's `zero_state` function
- The name "initial_state" is applied to the initial state.
- The `get_init_cell` function return the cell and initial state in the following tuple (Cell, InitialState)

Nice job stacking RNNs and initializing cell state.

To find out more about RNN effectiveness please follow this link:

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

The function `get_embed` applies embedding to `input_data` and returns embedded sequence.

Good job applying embedding to `input_data` and returning embedded sequence.

Check out this link:

https://www.tensorflow.org/api_docs/python/tf/contrib/layers/embed_sequence

to learn more about `tf.contrib.layers.embed_sequence()` which you could use to do the same.

The function `build_rnn` does the following:

- Builds the RNN using the `tf.nn.dynamic_rnn`.
- Applies the name "final_state" to the final state.
- Returns the outputs and final_state state in the following tuple (Outputs, FinalState)

The `build_nn` function does the following in order:

- Apply embedding to `input_data` using `get_embed` function.
- Build RNN using cell using `build_rnn` function.
- Apply a fully connected layer with a linear activation and `vocab_size` as the number of outputs.
- Return the logits and final state in the following tuple (Logits, FinalState)

The `get_batches` function create batches of input and targets using `int_text`. The batches should be a Numpy array of tuples. Each tuple is (batch of input, batch of target).

- The first element in the tuple is a single batch of input with the shape [batch size, sequence length]
- The second element in the tuple is a single batch of targets with the shape [batch size, sequence length]

Neural Network Training

- Enough epochs to get near a minimum in the training loss, no real upper limit on this. Just need to make sure the training loss is low and not improving much with more training.
- Batch size is large enough to train efficiently, but small enough to fit the data in memory. No real "best" value here, depends on GPU memory usually.
- Size of the RNN cells (number of units in the hidden layers) is large enough to fit the data well. Again, no real "best" value.
- The sequence length (`seq_length`) here should be about the size of the length of sentences you want to generate. Should match the structure of the data.
The learning rate shouldn't be too large because the training algorithm won't converge. But needs to be large enough that training doesn't take forever.

Set show every n batches to the number of batches the neural network should print progress

set show_every_n_batches to the number of batches the neural network should print progress.

Good job - your parameters look reasonable!

For more info about hyperparameter optimization please go [here](#)

The project gets a loss less than 1.0

Congratulations on getting the loss to 0.361

Generate TV Script

"input:0", "initial_state:0", "final_state:0", and "probs:0" are all returned by `get_tensor_by_name`, in that order, and in a tuple

The `pick_word` function predicts the next word correctly.

Great job using randomness when choosing the next word!

Your project predicts reasonable words!

The generated script looks similar to the TV script in the dataset.

It doesn't have to be grammatically correct or make sense.

Your generated script looks structurally similar to the TV script in the dataset! Congratulations!

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)

Rate this review

